

Outline

- **Part 1: Neural Style Transfer**
 - What is style transfer?
 - Style transfer implementation
 - Example: Generating images using style transfer
- **Part 2: Generative Adversarial Network (GAN)**
 - What is GAN?
 - GAN implementation
 - Training a GAN
- **Lab Assignment**

Part 1: Neural Style Transfer

Style Transfer

Style Source



Content Source



Output Image




Style Transfer Implementation

- Computing Loss
 - Content loss
 - Style loss
 - Total-variation regularization
- Put everything together: Style Transfer

Content Loss

- Content loss measures how much the feature map of the **generated image** differs from the feature map of the **source image**.
- We only care about the **content representation** of one layer of the network
- We will work with **reshaped versions** of these feature maps.



Content Loss

- $F^\ell \in \mathbb{R}^{C_\ell \times M_\ell}$: the feature map for the **current image**
- $P^\ell \in \mathbb{R}^{C_\ell \times M_\ell}$: the feature map for the **content source image**
- $M_\ell = H_\ell \times W_\ell = \#$ of elements in each feature map
- w_c = weight of the content loss term in the loss function

```
1 def content_loss(content_weight, content_current, content_original):
2     """
3     Compute the content loss for style transfer.
4
5     Inputs:
6     - content_weight: Scalar giving the weighting for the content loss.
7     - content_current: features of the current image; this is a PyTorch Tensor of shape
8       (1, C_l, H_l, W_l).
9     - content_target: features of the content image, Tensor with shape (1, C_l, H_l, W_l).
10
11     Returns:
12     - scalar content loss
13     """
14     loss = torch.sum( content_weight * (content_current - content_original)**2 )
15     return loss
```

The **content loss** is given by:

$$L_c = w_c \times \sum_{i,j} (F_{ij}^\ell - P_{ij}^\ell)^2$$


Style Transfer Implementation

- Computing Loss
 - Content loss
 - Style loss
 - Total-variation regularization
- Put everything together: Style Transfer

Style Loss

- For a given layer ℓ , the style loss is defined as follows:
 - First, compute the **Gram matrix G** which represents the correlations between the values in each channel of the feature map.
 - The **Gram matrix** is an approximation of the covariance matrix.
 - We want the **activation statistics** of our **generated image** to match the **activation statistics** of our **style image**.

Style Loss



The **feature map** is given by: $F^\ell \in \mathbb{R}^{C_\ell \times M_\ell}$

The **Gram matrix** is computed by: $G_{ij}^\ell = \sum_k F_{ik}^\ell F_{jk}^\ell$

The **Gram matrix** has shape: (C_ℓ, C_ℓ)

The **Gram matrix** from the feature map of the **current image**: G^ℓ

The **Gram matrix** from the feature map of the **source style image**: A^ℓ

Style Loss

The **style loss** for the layer ℓ is simply the weighted Euclidean distance between the two Gram matrices:

$$L_s^\ell = w_\ell \sum_{i,j} \left(G_{ij}^\ell - A_{ij}^\ell \right)^2$$

In practice we usually compute the style loss at a set of layers \mathcal{L} rather than just a single layer ℓ ; then the **total style loss** is the sum of style losses at each layer:

$$L_s = \sum_{\ell \in \mathcal{L}} L_s^\ell$$

```
1 def gram_matrix(features, normalize=True):
2     """
3     Compute the Gram matrix from features.
4
5     Inputs:
6     - features: PyTorch Tensor of shape (N, C, H, W) giving features for
7       a batch of N images.
8     - normalize: optional, whether to normalize the Gram matrix
9       If True, divide the Gram matrix by the number of neurons (H * W * C)
10
11     Returns:
12     - gram: PyTorch Tensor of shape (N, C, C) giving the
13       (optionally normalized) Gram matrices for the N input images.
14     """
15     N, C, H, W = features.size()
16     features = features.view(N*C,H*W) #reshape it
17     gram_matrix = torch.mm(features, features.t())
18     if (normalize):
19         gram_matrix /= float(H*W*C)
20     return gram_matrix
```

```
1 def style_loss(feats, style_layers, style_targets, style_weights):
2     """
3     Computes the style loss at a set of layers.
4     Inputs:
5     - feats: list of the features at every layer of the current image, as produced by
6       the extract_features function.
7     - style_layers: List of layer indices into feats giving the layers to include in the
8       style loss.
9     - style_targets: List of the same length as style_layers, where style_targets[i] is
10       a PyTorch Tensor giving the Gram matrix of the source style image computed at
11       layer style_layers[i].
12     - style_weights: List of the same length as style_layers, where style_weights[i]
13       is a scalar giving the weight for the style loss at layer style_layers[i].
14     Returns:
15     - style_loss: A PyTorch Tensor holding a scalar giving the style loss.
16     """
17     loss = 0
18     for i, layer in enumerate(style_layers):
19         current = gram_matrix(feats[layer])
20         loss += (style_weights[i] * torch.sum((current-style_targets[i])**2))
21     return loss
```

Style Transfer Implementation

- Computing Loss
 - Content loss
 - Style loss
 - Total-variation regularization
- Put everything together: Style Transfer

Total-variation regularization

- Encourage **smoothness** in the image by adding “**total variation**”
- Computed as the sum of the squares of differences in the pixel values for all pairs of pixels that are next to each other (**horizontally or vertically**)

```
1 def tv_loss(img, tv_weight):
2     """
3     Compute total variation loss.
4     Inputs:
5     - img: PyTorch Variable of shape (1, 3, H, W) holding an input image.
6     - tv_weight: Scalar giving the weight w_t to use for the TV loss.
7     Returns:
8     - loss: PyTorch Variable holding a scalar giving the total variation loss
9       for img weighted by tv_weight.
10    """
11    loss = 0
12    #do the row one
13    loss += torch.sum( ( img[:, :, 1:, :] - img[:, :, :-1, :] )**2 )
14    #on paper do for 2X2 and 3X3 then easy to see
15    #do the column one
16    loss += torch.sum( ( img[:, :, :, 1:] - img[:, :, :, :-1] )**2 )
17    #weighting
18    loss *= tv_weight
19    return loss
```

$$L_{tv} = w_t \times \left(\sum_{c=1}^3 \sum_{i=1}^{H-1} \sum_{j=1}^W (x_{i+1,j,c} - x_{i,j,c})^2 + \sum_{c=1}^3 \sum_{i=1}^H \sum_{j=1}^{W-1} (x_{i,j+1,c} - x_{i,j,c})^2 \right)$$

Style Transfer Implementation

- Computing Loss
 - Content loss
 - Style loss
 - Total-variation regularization
- Put everything together: Style Transfer

Put everything together

- Extract features for the content and style images separately
- Initialize output image to content or random noise
- Optimization setup
 - Hyperparameters
 - Optimizer choice
- Update loop
 - Compute losses
 - Update generated image

Example: Generating images with style transfer

Generating images with style transfer

- Use content image and style image to generate a new image
- Hyperparameters include:
 - The weights of content loss and style loss
 - The layer representations of content and style
 - Size of images

```
1 # Composition VII + Tubingen
2 params1 = {
3     'content_image' : 'tubingen.jpg',
4     'style_image' : 'composition_vii.jpg',
5     'image_size' : 192,
6     'style_size' : 512,
7     'content_layer' : 'conv_2',
8     'content_weight' : 5e-2,
9     'style_layers' : ['conv_1', 'conv_3', 'conv_5', 'conv_7'],
10    'style_weights' : (20000, 500, 1200, 1000),
11    'tv_weight' : 5e-2
12 }
13
14 style_transfer(**params1)
```


Generating images with style transfer



More weight to
content loss



More weight to
style loss

Generating images with style transfer

Content Reconstruction



Input Image



conv1_1



conv2_1



conv3_1



conv4_1




conv5_1

Style Reconstruction



Input Image



conv1_1



conv2_1



conv3_1



conv4_1



conv5_1

Generating images with style transfer

Resizing style image before running style transfer algorithm can transfer different types of features



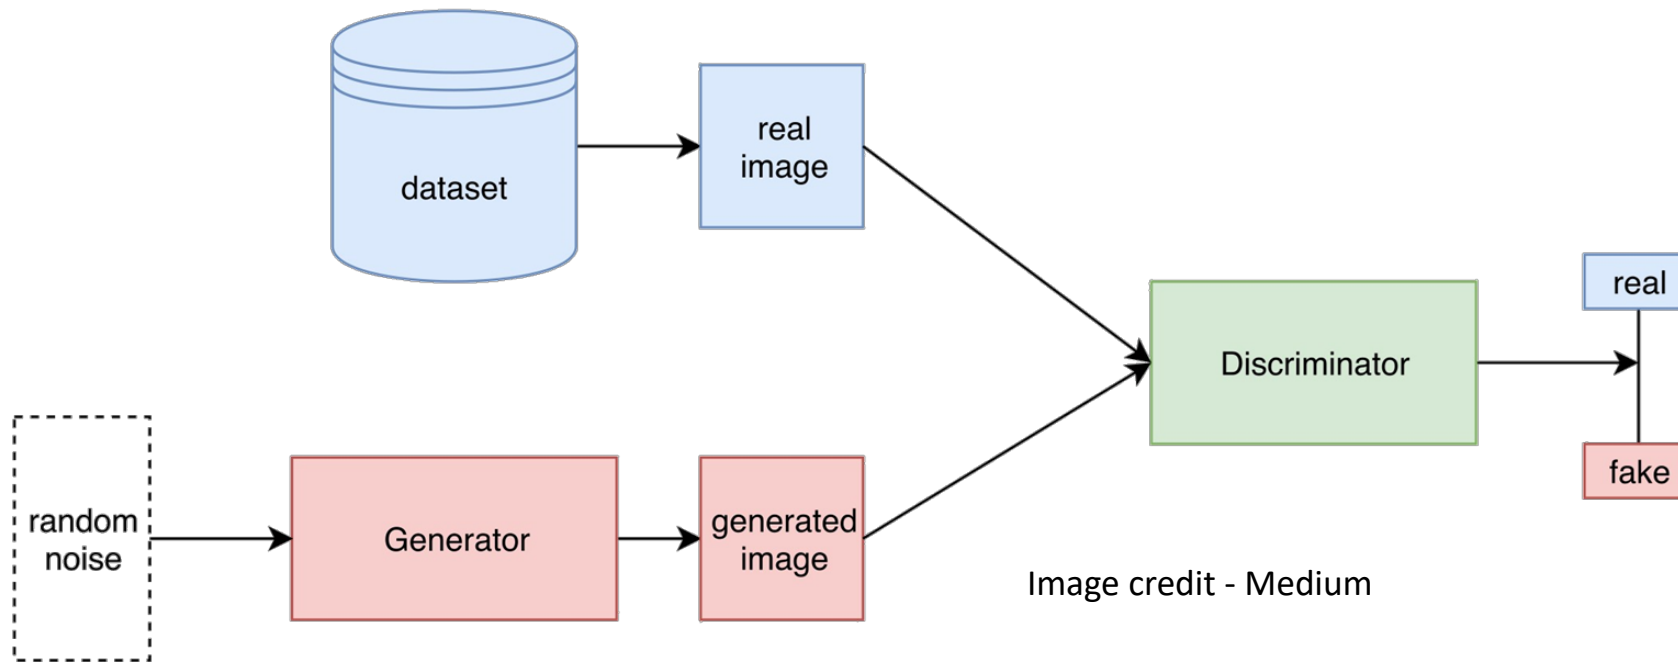
Larger style
image



Smaller style
image

Part 2: Generative Adversarial Networks (GANs)

What is a GAN?



What is a GAN?

- Cost Function

$$\underset{G}{\text{minimize}} \underset{D}{\text{maximize}} \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

- In practice, we alternate the following updates:

- Update the **generator** (G) to maximize the probability of the discriminator making the incorrect choice on generated data:


$$\underset{G}{\text{maximize}} \mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

- Update the **discriminator** (D), to maximize the probability of the discriminator making the correct choice on real and generated data:

$$\underset{D}{\text{maximize}} \mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] + \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

GAN Implementation

- Dataset: MNIST
- Random Noise
- Discriminator
- Generator
- GAN Loss
- Training a GAN



GAN Implementation

- Dataset: MNIST
- Random Noise
 - Generate uniform noise from -1 to 1 with shape [batch_size, dim].
 - You can also use torch.rand
- Discriminator
- Generator
- GAN Loss
- Training a GAN

```
1 def sample_noise(batch_size, dim):
2     """
3     Generate a PyTorch Tensor of uniform random noise.
4
5     Input:
6     - batch_size: Integer giving the batch size of noise to generate.
7     - dim: Integer giving the dimension of noise to generate.
8
9     Output:
10    - A PyTorch Tensor of shape (batch_size, dim) containing uniform
11      random noise in the range (-1, 1).
12    """
13    return torch.FloatTensor(batch_size, dim).uniform_(-1, 1)
```


GAN Implementation

- Dataset: MNIST
- Random Noise
- Discriminator
 - FC layer (784 -> 256)
 - LeakyRelu with alpha 0.01
 - FC layer (256 -> 256)
 - LeakyRelu with alpha 0.01
 - FC layer (256 -> 1)
- Generator
- GAN Loss
- Training a GAN

```
1 def discriminator():
2     """
3     Build and return a PyTorch model implementing the architecture above.
4     """
5     model = nn.Sequential( Flatten(),
6                             nn.Linear(784, 256),
7                             nn.LeakyReLU(inplace=True),
8                             nn.Linear(256, 256),
9                             nn.LeakyReLU(inplace=True),
10                            nn.Linear(256, 1)
11                        )
12     return model
```

GAN Implementation

- Dataset: MNIST
- Random Noise
- Discriminator
- Generator
 - FC layer (noise_dim -> 1024)
 - Relu
 - FC layer(1024 -> 1024)
 - Relu
 - FC layer (1024 -> 784)
 - Tanh (to clip the image to be in the range [-1, 1])
- GAN Loss
- Training a GAN

```
1 def generator(noise_dim=NOISE_DIM):
2     """
3     Build and return a PyTorch model implementing the architecture above.
4     """
5     model = nn.Sequential( nn.Linear(noise_dim,1024),
6                           nn.ReLU(inplace=True),
7                           nn.Linear(1024,1024),
8                           nn.ReLU(inplace=True),
9                           nn.Linear(1024,784),
10                          nn.Tanh()
11                      )
12     return model
```

GAN Implementation

- Dataset: MNIST
- Random Noise
- Discriminator
- Generator
- **GAN Loss**
- Training a GAN

GAN Loss

- The generator loss:

$$\ell_G = -\mathbb{E}_{z \sim p(z)} [\log D(G(z))]$$

- The discriminator loss:

$$\ell_D = -\mathbb{E}_{x \sim p_{\text{data}}} [\log D(x)] - \mathbb{E}_{z \sim p(z)} [\log(1 - D(G(z)))]$$

- We will use the binary cross entropy loss - Given a score s and a label y , the binary cross entropy loss is:

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

GAN loss

- A naïve implementation of the binary cross entropy loss formula can be numerically unstable.

$$bce(s, y) = -y * \log(s) - (1 - y) * \log(1 - s)$$

```
1 def bce_loss(input, target):
2     """
3     Numerically stable version of the binary cross-entropy loss function.
4
5     As per https://github.com/pytorch/pytorch/issues/751
6     See the TensorFlow docs for a derivation of this formula:
7     https://www.tensorflow.org/api\_docs/python/tf/nn/softmax\_cross\_entropy\_with\_logits
8
9     Inputs:
10    - input: PyTorch Tensor of shape (N, ) giving scores.
11    - target: PyTorch Tensor of shape (N,) containing 0 and 1 giving targets.
12
13    Returns:
14    - A PyTorch Tensor containing the mean BCE loss over the minibatch of input data.
15    """
16    neg_abs = - input.abs()
17    loss = input.clamp(min=0) - input * target + (1 + neg_abs.exp()).log()
18    return loss.mean()
```

GAN Loss

- Discriminator loss implementation

```
1 def discriminator_loss(logits_real, logits_fake):
2     """
3     Computes the discriminator loss described above.
4     Inputs:
5     - logits_real: PyTorch Tensor of shape (N,) giving scores for the real data.
6     - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.
7     Returns:
8     - loss: PyTorch Tensor containing (scalar) the loss for the discriminator.
9     """
10    N, _ = logits_real.size()
11    loss = (bce_loss(logits_real, torch.ones(N).type(dtype)))+(bce_loss(logits_fake, torch.zeros(N).type(dtype)))
12    return loss
```

- Generator loss implementation

```
1 def generator_loss(logits_fake):
2     """
3     Computes the generator loss described above.
4     Inputs:
5     - logits_fake: PyTorch Tensor of shape (N,) giving scores for the fake data.
6     Returns:
7     - loss: PyTorch Tensor containing the (scalar) loss for the generator.
8     """
9     N, _ = logits_fake.size()
10    loss = (bce_loss(logits_fake, torch.ones(N).type(dtype)))
11    return loss
```

Training a GAN

- **Training Discriminator**

- Pass the **sample noise** to generator and get the fake images
- Use **.detach()** to break gradient connection
- Pass the fake images to discriminator, compute loss, and update discriminator

- **Training Generator**

- Pass the sample noise to generator and get the fake images
- Pass the fake images to discriminator
- Compute generator loss and update generator


```
1 def run_a_gan(D, G, D_solver, G_solver, discriminator_loss, generator_loss, show_every=250,
2             batch_size=128, noise_size=96, num_epochs=10):
3     """
4     Train a GAN!
5
6     Inputs:
7     - D, G: PyTorch models for the discriminator and generator
8     - D_solver, G_solver: torch.optim Optimizers to use for training the
9       discriminator and generator.
10    - discriminator_loss, generator_loss: Functions to use for computing the generator and
11      discriminator loss, respectively.
12    - show_every: Show samples after every show_every iterations.
13    - batch_size: Batch size to use for training.
14    - noise_size: Dimension of the noise to use as input to the generator.
15    - num_epochs: Number of epochs over the training dataset to use for training.
16    """
17    iter_count = 0
18    for epoch in range(num_epochs):
19        for x, _ in loader_train:
20            if len(x) != batch_size:
21                continue
22            D_solver.zero_grad()
23            real_data = x.type(dtype)
24            logits_real = D(2* (real_data - 0.5)).type(dtype)
25
26            g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
27            fake_images = G(g_fake_seed).detach()
28            logits_fake = D(fake_images.view(batch_size, 1, 28, 28))
29
30            d_total_error = discriminator_loss(logits_real, logits_fake)
31            d_total_error.backward()
32            D_solver.step()
33
34            G_solver.zero_grad()
35            g_fake_seed = sample_noise(batch_size, noise_size).type(dtype)
36            fake_images = G(g_fake_seed)
37
38            gen_logits_fake = D(fake_images.view(batch_size, 1, 28, 28))
39            g_error = generator_loss(gen_logits_fake)
40            g_error.backward()
41            G_solver.step()
```

Training a GAN


```
1 # Make the discriminator
2 D = discriminator().type(dtype)
3
4 # Make the generator
5 G = generator().type(dtype)
6
7 # Use the function you wrote earlier to get optimizers for the Discriminator and the Generator
8 D_solver = get_optimizer(D)
9 G_solver = get_optimizer(G)
10 # Run it!
11 run_a_gan(D, G, D_solver, G_solver, discriminator_loss, generator_loss)
```



Iteration 0



Iteration 1000



Iteration 2000



Iteration 3000


Lab Assignment: Deep Convolutional GANs

Assignment Details

- Implement a deep convolutional GAN (DCGAN) to generate much better fake images
- Implement the **discriminator** with:
 - Conv2D: 32 filters, 5x5, Stride 1
 - Leaky Relu (alpha = 0.01)
 - Max Pool 2x2, Stride 2
 - Conv2D: 64 filters, 5x5, Stride 1
 - Leaky ReLu (alpha = 0.01)
 - Max Pool 2x2 Stride 2
 - Flatten
 - Fully Connected with output size 4x4x64
 - Leaky Relu (alpha = 0.01)
 - Fully Connected with output size 1



Results of GAN




Results of DCGAN

Assignment Details

- Implement the **generator** with:
 - Fully connected with output size 1024
 - Relu
 - BatchNorm
 - Fully connected with output size 7x7x128
 - ReLu
 - BatchNorm
 - Reshape into image Tensor of shape 7,7,128
 - Conv2D Transpose: 64 filters of 4x4, stride 2, 'same' padding (use padding = 1)
 - BatchNorm
 - Conv2D Transpose: 1 filter of 4x4, stride 2, 'same' padding (use padding = 1)
 - Tanh
 - Should have a 28x28x1 image, reshape back into 784 vector



Results of GAN



Results of DCGAN

Evaluation:

Train DCGAN on MNIST dataset and plot out the generated images every 150 training iterations for 5 epochs